

Docker Security

An analysis of security threats and recommended practices for building a secure Docker infrastructure

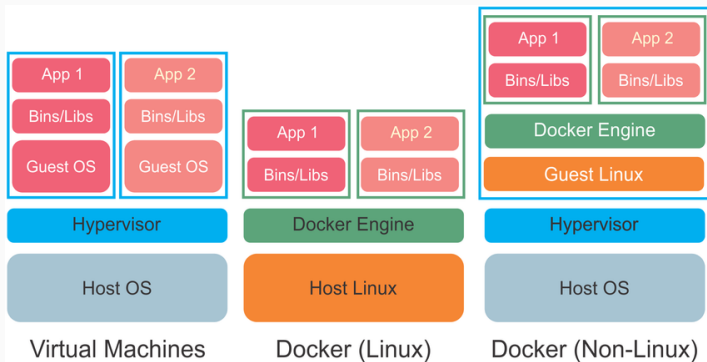
Andreas Pfefferle

August 24, 2018

1. Introduction to Docker
2. Built-in Security Features of Docker
3. Docker Security Threats
4. Recommendations for Security Improvements

Introduction to Docker

Comparison with Hypervisor-based Virtualization

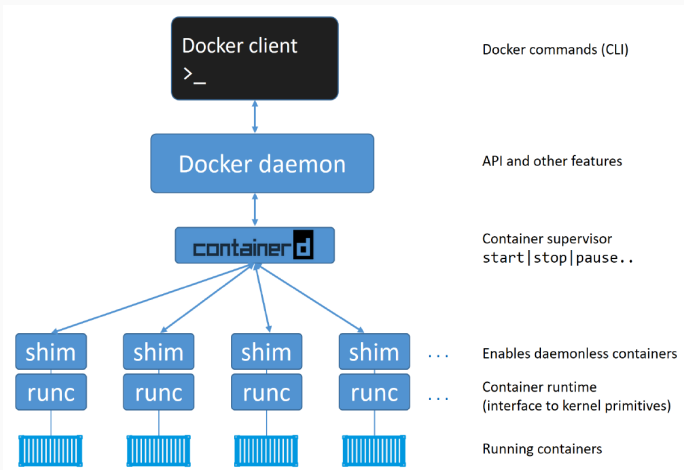


(Docker Inc. 2013)

History


- 2000: FreeBSD *jails*
- 2002-2008: namespaces, capabilities & control groups were added to the Linux kernel
- 2004: Solaris *Zones*
- 2005: Open VZ
- 2008: Linux Containers (LXC)
- 2013: Docker (with LXC)
- 2014: Docker's self-developed *libcontainer* replaced LXC in Docker
- 2015: Open Container Initiative
- 2015: Docker *runC*
- 2016: Docker *containerd*
- 2017: Moby

Architecture



Docker Engine (Poulton 2016)

Images



Docker filesystem layers

```
FROM ubuntu:18.04
LABEL John Doe "john.doe@johndoe.com"
RUN apt-get update; apt-get install -y nginx
RUN echo 'Hello world!' \
    >/var/www/html/index.html
VOLUME /app
EXPOSE 80
ENTRYPOINT ["/usr/sbin/nginx", "-g", "daemon off;"]
```

Dockerfile for a simple web server

Image Build and Container Start

```
---> c5984a6c0671
Step 4/7 : RUN echo 'Hello world!' >/var/www/html/index.html
---> Running in 933d0c1dc42c
Removing intermediate container 933d0c1dc42c
---> dd8b292f775d
Step 5/7 : VOLUME /app
---> Running in 3f9c4114ca4e
Removing intermediate container 3f9c4114ca4e
---> 04dd9c606b31
Step 6/7 : EXPOSE 80
---> Running in f008bb503ccd
Removing intermediate container f008bb503ccd
---> e74f0cb10744
Step 7/7 : ENTRYPOINT ["/usr/sbin/nginx", "-g", "daemon off;"]
---> Running in b23a2b6f4382
Removing intermediate container b23a2b6f4382
---> a26689f3704c
Successfully built a26689f3704c
Successfully tagged secomba/webserver:latest
```

Running `$ sudo docker build -t="secomba/webserver:latest" .`


Image Build and Container Start


```
$ sudo docker run -t -i secomba/webserver:latest
```

Image Build and Container Start

```
andi@callisto:~$ sudo docker run -t -i dd8b292f775d /bin/bash
root@7e7d3ff517ee:/# ls
bin boot dev etc home lib lib64 media mnt opt proc root sbin sys
root@7e7d3ff517ee:/# cat /var/www/html/index.html
Hello world!
root@7e7d3ff517ee:/#
```





Docker Registries



Dashboard Explore Organizations Create  anpsecomba

Repositories (48759)

All

	node official	6.0K STARS	10M+ PULLS	> DETAILS
	nodered/node-red-docker public automated build	174 STARS	10M+ PULLS	> DETAILS
	bitnami/node public automated build	30 STARS	10M+ PULLS	> DETAILS
	apptsvc/node public automated build	11 STARS	10M+ PULLS	> DETAILS

Docker Hub

Built-in Security Features of Docker

Linux Namespaces for Container Isolation

- Process isolation through the PID namespace
- Network Isolation through the network namespace
- File system isolation through the mount namespace
- Isolation of inter-process communication through the IPC namespace
- UTS isolation through the UTS namespace
- (User isolation through the user namespace)
- ~~Control Group isolation through the cgroup namespace~~

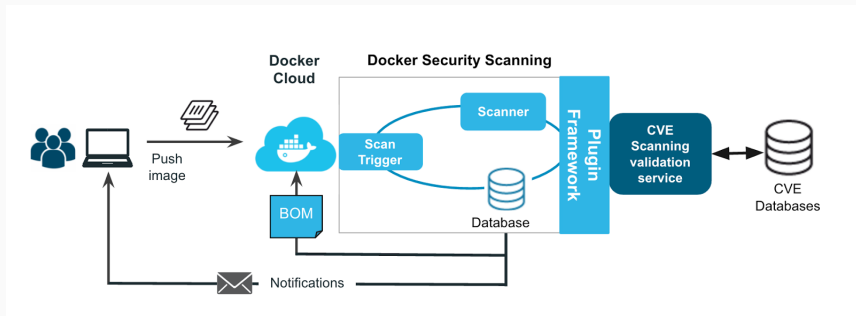
Linux Capabilities to restrict Access Rights

- Access Control in Unix: Discretionary Access Control (DAC)
- Linux capabilities: Dividing root privileges into more than 30 individual capabilities
- Docker: currently 14 capabilities
- Docker run command allows `cap-add` and `cap-drop` parameters

Resource Management through Control Groups

- Control groups manage the host's resources
- Docker: individually allocate resources to each container
- Control groups try to prevent Denial of Service attacks

Docker Security Scanning for Image Inspection



Docker Security Scanning

(Toli Kuznets 2016)

Docker Security Scanning for Image Inspection

10.7.0-alpine Compressed size: 25 MB

Scanned 2 days ago

! This image has vulnerabilities



10 Compressed size: 266 MB

Scanned 2 days ago

! This image has vulnerabilities



10.7 Compressed size: 266 MB

Scanned 2 days ago

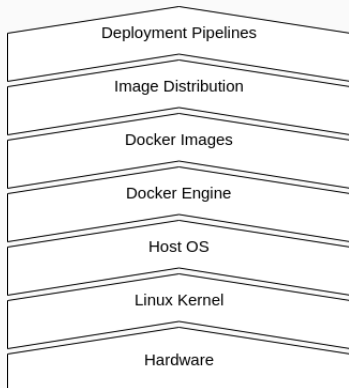
! This image has vulnerabilities



Multiple images in the official Node repository

Docker Security Threats

Overview and Approach



Overview of the investigated layers. Each layer will be analyzed with STRIDE to identify the threats.

Identified Threats:

Spoofing -

Tampering -

Repudiation -

Information Disclosure Unauthorized processes can read
arbitrary memory

Denial of Service -

Elevation of Privilege -

Exploitation:

- Meltdown (Kernel patches mitigated the risk)
- Spectre (multiple variants disclosed regularly)
- Prime+Probe in Intel SGX (worked in a Docker environment)

Identified Threats:

Spoofing -

Tampering Using a memory corruption vulnerability in the kernel, arbitrary code can be executed. Thus, an attacker could modify files, which violates the integrity of the files.

Repudiation -

Information Disclosure -

Denial of Service -

Elevation of Privilege -

Exploitation:

- Missing pointer checks lead to memory corruption
 - Attacker uses a kernel function and passes a pointer value
 - Kernel should use the `access_ok` function
 - Kernel does not check whether the value of the provided pointer points to only user-space memory
 - Attacker can read and write to arbitrary kernel locations and execute arbitrary code

Identified Threats:

Spoofing -

Tampering -

Repudiation -

Information Disclosure -

Denial of Service -

Elevation of Privilege Mounting the host's root filesystem with write permissions enables the execution of highly privileged executable program files.

Exploitation:

- Mounting `/var/run/docker.sock` into a container: e.g.,
`nginx-proxy`

Identified Threats:

Spoofing -

Tampering As in CVE-2014-6408, containers can bypass the isolation. It has been shown that containers obtain read and write access on the entire host file system in the event of an outbreak.

Repudiation -

Information Disclosure -

Denial of Service -

Elevation of Privilege -

Exploitation:

- 3 CVE vulnerabilities have not yet been mitigated

Identified Threats:

Spoofing -

Tampering -

Repudiation -

Information Disclosure -

Denial of Service A maliciously prepared image from a public repository uses the computing power of the machine to mine cryptocurrencies.

Elevation of Privilege -

Exploitation:

- Docker Hub repository "docker123321": 5 million downloads, around 90,000\$ Monero coins

Identified Threats:

Spoofing -

Tampering -

Repudiation -

Information Disclosure -

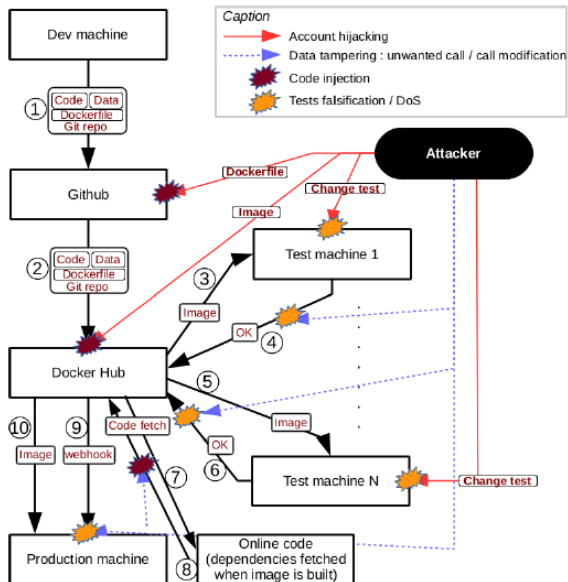
Denial of Service Unpacking untrusted zipped archives may cost a lot of computing power and, eventually, lead to full memory.

Elevation of Privilege -

Exploitation:

- Zip Bomb: 42.zip

Deployment Pipelines



Exploitation:

- 5 minutes and 30 seconds after the push to Github: container with malicious image runs in production machine
- July 2018: 22,000 management interfaces publicly accessible, 300 without a password

Recommendations for Security Improvements

- Intel Cascade Lake processors in the second half of 2018

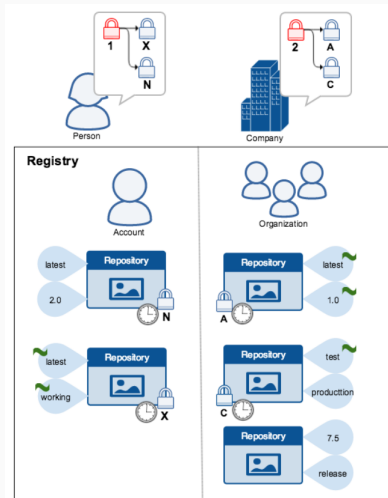
- Updating the kernel regularly
- Unikernels?

- Using a host operating system that supports Mandatory Access Control (MAC) such as AppArmor or SELinux
- Seccomp: filter incoming system calls using a whitelist
- Removing all capabilities except those that are explicitly required
- Reducing mount operations

- User namespace: root privileges inside the container, but is effectively mapped to an unprivileged UID on the host system

- If a service does not need root privileges inside the container:
Dockerfile USER instruction
- Smaller base images = smaller attack surface
- Remove package installer
- Local security scans with Anchore or Clair
- Avoid vulnerability-prone technologies

Image Distribution



Docker Content Trust

- 2FA, U2F (e.g., Google: no successful phishing attack in over a year and a half)
- fewer automation in automated pipelines

Conclusion

- complexity overload, especially with Kubernetes, Docker Swarm, etc.
- only on dedicated hardware
- only with hardening measures